

Offline Operational Report

This document was generated 100% offline using INDRA's unified document suite.

Executive Summary

All components are functioning within nominal parameters. Local inferencing, offline tool execution, and workspace persistence are operational.

System Performance Metrics

Metric	Baseline	Current	Status
Inference Latency	45ms	38ms	Optimal
VRAM Footprint	5.8 GB	5.2 GB	Protected
PDF Generation Speed	1.8s	0.4s	Accelerated
Token Throughput	24 tps	29 tps	Exceeded

Key Recommendations

- **Maintain local mode:** Keep ``approvals.mode: smart`` for user oversight or ``off`` for autonomous execution.
- **Utilize Platypus flowables:** Prevents text clipping and maintains page boundaries.
- **Export directly to `./outputs/`:** Preserves deliverables across session teardowns.